

Capacity Building in Seasonal Hydrological Forecasting

Modeling and Machine Learning

AGRHYMET, Climate Regional Center for West-Africa and Sahel

@Arsène KIEMA

2025-10-06



Pedagogical Objectives

Learning outcomes

By the end of this module, participants will be able to:

- Understand the theoretical foundations of four key models: PCR, Ridge, Lasso, and Random Forest.
- Explain how these models handle collinearity, regularization, and overfitting.
- Train and evaluate each model in R using hydrological data.
- Interpret model coefficients, variable importance, and performance metrics.

Practical Implementation in R

We now apply these models to a simple rainfall–runoff dataset.

Setting Up the Environment

```
library(tidyverse)
library(tidymodels)
library(ranger)
library(rsample)
library(lubridate)
library(hydroGOF)
```

Practical Implementation in R

Load a sample hydro-climatic dataset

```
data <- read.csv("data/hydro_features.csv") |>  
  mutate(Date = as.Date(Date)) |>  
  arrange(Date) |>  
  drop_na()
```

Practical Implementation in R

Define rolling-origin resamples

We define training and testing periods that move forward in time.

```
n_initial <- ceiling(nrow(data) * 0.7)
n_assess  <- floor(nrow(data) * 0.2)
n_skip    <- 1

ro <- rolling_origin(
  data,
  initial      = n_initial,
  assess       = n_assess,
  skip         = n_skip,
  cumulative   = TRUE
)
```

Practical Implementation in R

Recipe design to avoid leakage

All preprocessing (centering, scaling, PCA, etc.) is done **inside** the recipe. `fit_resamples()` and `tune_grid()` will re-estimate these steps **within each split**.

Principal Component Regression (PCR)

PCR reduces correlated predictors into principal components before regression. This helps when rainfall, PET, and temperature are strongly related.

```
rec_pcr <- recipe(Qobs ~ Rain + Temp + PET, data = data) |>
  step_zv(all_predictors()) |>
  step_normalize(all_predictors()) |>
  step_pca(all_predictors(), num_comp = tune())

mod_pcr <- linear_reg() |> set_engine("lm")

wf_pcr <- workflow() |> add_recipe(rec_pcr) |> add_model(mod_pcr)

grid_pcr <- tibble(num_comp = 1:3)
```

Principal Component Regression (PCR)

Train PCR

Train and evaluate the PCR model across all time splits:

```
set.seed(123)
res_pcr <- tune_grid(
  wf_pcr,
  resamples = ro,
  grid = grid_pcr,
  metrics = metric_set(yardstick::rmse, yardstick::rsq),
  control = control_grid(save_pred = TRUE)
)
show_best(res_pcr, metric = "rmse")

# A tibble: 3 x 7
  num_comp .metric .estimator mean      n std_err .config
  <int> <chr> <chr> <dbl> <int> <dbl> <chr>
```


Principal Component Regression (PCR)

Train PCR

```
# A tibble: 3 x 7
```

	num_comp	.metric	.estimator	mean	n	std_err	.config
	<int>	<chr>	<chr>	<dbl>	<int>	<dbl>	<chr>
1	3	rmse	standard	549.	3	20.1	pre3_mod0_post0
2	2	rmse	standard	610.	3	43.2	pre2_mod0_post0
3	1	rmse	standard	611.	3	36.9	pre1_mod0_post0

Ridge Regression (L2 Regularization)

Ridge regression reduces overfitting by shrinking large coefficients.

```
mod_ridge <- linear_reg(penalty = tune(), mixture = 0) |> set_engine("glmnet")

rec_ridge <- recipe(Qobs ~ Rain + Temp + PET, data = data) |>
  step_normalize(all_predictors())

wf_ridge <- workflow() |> add_model(mod_ridge) |> add_recipe(rec_ridge)
```

Ridge Regression (L2 Regularization)

Train Ridge

```
set.seed(123)
res_ridge <- tune_grid(
  wf_ridge,
  resamples = ro,
  grid = 20,
  metrics = metric_set(yardstick::rmse, yardstick::rsq),
  control = control_grid(save_pred = TRUE)
)

show_best(res_ridge, metric = "rmse")
```

Ridge Regression (L2 Regularization)

Train Ridge

```
# A tibble: 5 x 7
```

	penalty	.metric	.estimator	mean	n	std_err	.config
	<dbl>	<chr>	<chr>	<dbl>	<int>	<dbl>	<chr>
1	1.21e-10	rmse	standard	549.	3	20.7	pre0_mod01_post0
2	3.71e-10	rmse	standard	549.	3	20.7	pre0_mod02_post0
3	1.00e- 9	rmse	standard	549.	3	20.7	pre0_mod03_post0
4	3.23e- 9	rmse	standard	549.	3	20.7	pre0_mod04_post0
5	1.07e- 8	rmse	standard	549.	3	20.7	pre0_mod05_post0

Lasso Regression (L1 Regularization)

Lasso penalizes coefficients and sets some of them to zero (variable selection).

```
mod_lasso <- linear_reg(penalty = tune(), mixture = 1) |> set_engine(  
  
rec_lasso <- recipe(Qobs ~ Rain + Temp + PET, data = data) |>  
  step_normalize(all_predictors())  
  
wf_lasso <- workflow() |> add_model(mod_lasso) |> add_recipe(rec_lasso)
```

Lasso Regression (L1 Regularization)

Train Lasso

```
set.seed(123)
res_lasso <- tune_grid(
  wf_lasso,
  resamples = ro,
  grid = 20,
  metrics = metric_set(yardstick::rmse, yardstick::rsq),
  control = control_grid(save_pred = TRUE)
)

show_best(res_lasso, metric = "rmse")
```

Lasso Regression (L1 Regularization)

Train Lasso

```
# A tibble: 5 x 7
```

	penalty	.metric	.estimator	mean	n	std_err	.config
	<dbl>	<chr>	<chr>	<dbl>	<int>	<dbl>	<chr>
1	1.21e-10	rmse	standard	549.	3	20.2	pre0_mod01_post0
2	3.71e-10	rmse	standard	549.	3	20.2	pre0_mod02_post0
3	1.00e- 9	rmse	standard	549.	3	20.2	pre0_mod03_post0
4	3.23e- 9	rmse	standard	549.	3	20.2	pre0_mod04_post0
5	1.07e- 8	rmse	standard	549.	3	20.2	pre0_mod05_post0

Random Forest

Random Forest is an ensemble of trees, ideal for nonlinear rainfall–runoff relationships.

```
mod_rf <- rand_forest(  
  mtry = tune(),  
  min_n = tune(),  
  trees = 500  
) |>  
  set_engine("ranger", importance = "permutation") |>  
  set_mode("regression")  
  
rec_rf <- recipe(Qobs ~ Rain + Temp + PET, data = data)  
  
wf_rf <- workflow() |> add_model(mod_rf) |> add_recipe(rec_rf)
```


Random Forest

Training

```
set.seed(123)
res_rf <- tune_grid(
  wf_rf,
  resamples = ro,
  grid = grid_rf,
  metrics = metric_set(yardstick::rmse, yardstick::rsq),
  control = control_grid(save_pred = TRUE)
)

show_best(res_rf, metric = "rmse")
```

Random Forest

Training

```
# A tibble: 5 x 8
```

	mtry	min_n	.metric	.estimator	mean	n	std_err	.config
	<int>	<int>	<chr>	<chr>	<dbl>	<int>	<dbl>	<chr>
1	3	15	rmse	standard	572.	3	57.6	pre0_mod14_pos
2	2	11	rmse	standard	572.	3	54.4	pre0_mod08_pos
3	2	15	rmse	standard	574.	3	55.7	pre0_mod09_pos
4	1	2	rmse	standard	575.	3	53.7	pre0_mod01_pos
5	3	20	rmse	standard	575.	3	53.3	pre0_mod15_pos

Comparing Model Performances

We aggregate results from all models to identify the best one.

```
cv_tbl <- bind_rows(  
  mutate(collect_metrics(res_pcr), model = "PCR"),  
  mutate(collect_metrics(res_ridge), model = "Ridge"),  
  mutate(collect_metrics(res_lasso), model = "Lasso"),  
  mutate(collect_metrics(res_rf), model = "RandomForest")  
)
```

Comparing Model Performances

We aggregate results from all models to identify the best one.

```
# A tibble: 6 x 11
```

	num_comp	.metric	.estimator	mean	n	std_err	.config	model
	<int>	<chr>	<chr>	<dbl>	<int>	<dbl>	<chr>	<chr>
1	1	rmse	standard	6.11e+2	3	3.69e+1	pre1_m~	PCR
2	1	rsq	standard	5.17e-3	3	4.45e-3	pre1_m~	PCR
3	2	rmse	standard	6.10e+2	3	4.32e+1	pre2_m~	PCR
4	2	rsq	standard	5.47e-3	3	5.00e-3	pre2_m~	PCR
5	3	rmse	standard	5.49e+2	3	2.01e+1	pre3_m~	PCR
6	3	rsq	standard	4.89e-2	3	2.58e-2	pre3_m~	PCR

```
# i 1 more variable: min_n <int>
```

**THANK YOU FOR YOUR
ATTENTION**

